Companion — Full Pipeline Report (v3)

Optum / UHG · updated skill (theme-level Prioritization, PRD Technical Considerations, RICE-ranked backlog) · includes illustrative PMF scorecard

WHAT CHANGED FROM V1

Stage 8 now scores 5 themes instead of 10 features, cascading a genuinely different decomposition through Stages 9–13 — Theme B ("Agent Handoff Quality") split into two real features on proper decomposition. Stage 10's PRD now includes a Technical Considerations section (NFRs, integration approach, rollout/rollback). Stage 14's backlog is RICE-ranked instead of Status-ordered. Stage 15 adds an illustrative best-judgment scorecard alongside the real verdict. Stages 1–7, 12 are otherwise unchanged.

CONTENTS

PHASE: DISCOVERY — UNCHANGED

Stage 1 — Discovery

Initial read: Optum's contact center treats all inquiries — from simple deductible questions to genuinely complex claims disputes — as equally deserving of a live agent. That's expensive (~$57M/yr) and slow, and it's specifically the routine, repetitive segment of volume that's misallocated. Companion is scoped as a triage-and-resolve layer, not a replacement for agents: absorb the easy tail, hand off the hard cases with context.

Main Hypotheses

- H1: A large share of contact volume is answerable from data UHG already has, without human judgment.
- H2: Members abandon self-service today mainly because past bots were generic/non-personalized, not because they inherently prefer humans.
- H3: Deep systems integration + HIPAA-compliant personalization is the actual differentiator.
- H4: Warm transfer with context reduces AHT on escalated calls even when containment misses target.
- H5: Providers and patients have different enough intents/systems that a single undifferentiated bot flow will underperform for one or both groups.

Open Questions (never fully resolved — carried through the pipeline)

- Volume mix / query taxonomy by category — never obtained; still blocks the taxonomy work item.
- Actual integration lift for claims/benefits/provider-directory — weeks vs. quarters.
- Whether the $5M→$30M savings model is bottom-up or a top-down target.

Assumptions Used (best-judgment, flagged)

- The $57M is predominantly labor/AHT-driven contact-center spend.
- Roughly 60–70% of inbound volume is "routine."
- Patients and providers are two distinct segments with different systems/needs.
- Current self-service resolution rate is effectively ~0%; CSAT baseline is unmeasured pre-MVP.

PHASE: DISCOVERY — UNCHANGED

Stage 2 — Persona Mapping

Hypothesis-based, audience: PM/eng requirements. Three personas — one per segment: Patient, Provider, internal Support Agent.

Companion — Stage 2 personas

Hypothesis-based · no formal research yet · PM/eng requirements audience

Patient — the routine-question member

Hypothesis-based

UHG member contacting Optum about benefits, coverage, claims, or billing

GOALS

Fast, correct answer without holding
No re-explaining their situation
Trust the answer reflects their real account

PAIN POINTS

Long wait times for simple questions
Past bad experiences with generic bots
Doubt about whether a bot "really knows" them

> *"I just want to know if this is covered without sitting on hold for 20 minutes." [HYPOTHESIS: unvalidated]*

Validation status: fully hypothesis-based. Upgrade with member interviews or ticket-category data.

Provider — the time-pressed verifier

Hypothesis-based

Provider or office staff checking benefits/coverage mid-workflow

GOALS

Verify coverage/eligibility quickly
Minimize hold time blocking workflow
Get an answer they can act on now

PAIN POINTS

Wait time costs more mid-appointment
May need to check for an absent patient
No good "call back later" option

> *"I need this answer in the next two minutes, not after being on hold." [HYPOTHESIS: unvalidated]*

Validation status: fully hypothesis-based. Upgrade with provider-office interviews or provider-channel ticket data.

Support agent — the escalation receiver

Hypothesis-based · internal

Internal / operational segment — workflow-centered, not consumer-shaped

GOALS

Resolve escalations fast without re-collecting info
Trust the handed-off context is accurate

PAIN POINTS

No context passed on transfer today
Risk of noisy or wrong bot handoffs adding cleanup work

Validation status: fully hypothesis-based. Upgrade with agent interviews or shadowing.

PHASE: DISCOVERY — UNCHANGED

Stage 3 — Empathy Map

Design implications: Patient — surface account context early to counter distrust. Provider — support patient-absent lookups, terse tone. Support agent — warm transfer must deliver structured context before the call connects.

Companion — Stage 3 empathy maps

Hypothesis-based · derived from Stage 2 personas only · informs interaction/tone/escalation design

Patient — the routine-question member

See Persona: routine-question member

SAYS "Is this covered?" / "Why was this denied?" / "I don't want to explain this again."	**THINKS** Wonders if the bot actually knows their account, or if it's another generic dead end *[INFERRED: from stated distrust of past bots]*
FEELS Frustrated by the wait; *[INFERRED: mildly anxious about cost/coverage uncertainty]*	**DOES** Calls or uses the member portal today; likely abandons self-service if it feels generic

PAINS

Long hold times for simple questions; burned by non-personalized bots before; unsure the bot "knows" them

GAINS

A fast, correct answer grounded in their actual plan/claim data, without repeating their situation

DESIGN IMPLICATION

Companion should surface specific account context early in the conversation (e.g., "I can see your claim from [date]") to signal it isn't generic — this directly counters the stated distrust that's the core adoption barrier for this persona.

Provider — the time-pressed verifier

See Persona: time-pressed verifier

SAYS "Is this covered for this patient?" / "I need this now, not a callback."	**THINKS** Weighs whether self-service will actually handle a patient-absent lookup, or if it'll dead-end into a hold anyway *[INFERRED: from stated pain about absent-patient checks]*
FEELS *[INFERRED: time pressure/urgency]* — goal is speed, pain is a blocked workflow	**DOES** Uses the provider portal and phone in parallel today

PAINS

Wait time blocks clinical/office workflow; needs to check coverage for a patient who isn't present; no usable "call back later" option

GAINS

A fast, actionable coverage answer they can act on immediately, without switching channels

DESIGN IMPLICATION

Companion should support provider-initiated, patient-absent coverage lookups directly, and default to terse, high-density responses for provider sessions rather than the more conversational tone appropriate for patients — urgency, not warmth, is the dominant need here.

Support agent — the escalation receiver

See Persona: escalation receiver (internal)

SAYS "What has the member already been asked / told?" *[INFERRED: implied by the stated pain of no context today]*	**THINKS** Concerned they'll have to redo the member's request from scratch if no context arrives *[INFERRED]*
FEELS *[INFERRED: mild skepticism]* toward bot handoffs, given no context is passed today	**DOES** Works from whatever the member says live, plus internal systems — no pre-bot transcript today

PAINS

No context passed on transfer today; risk that bot handoffs are noisy or wrong, adding cleanup work

GAINS

A complete, accurate context summary handed off automatically — less repeated questioning, lower AHT

DESIGN IMPLICATION

Warm transfer must deliver a structured summary (member intent, data already surfaced, why the bot couldn't resolve it) to the agent's screen before the call connects — not just route the call — since agent trust in handoff quality is the real blocker to any AHT improvement.

PHASE: DISCOVERY — UNCHANGED

Stage 4 — Customer Journey Map

Current-state maps for all three personas. Support Agent's "Close" stage flagged as a real gap rather than padded.

Companion — Stage 4 journey maps (current-state)

Hypothesis-based · derived from Stage 2/3 personas and empathy maps only · informing UX design

Patient — the routine-question member

Current-state · stages: Trigger → Attempt Self-Service → Wait/Escalate → Get Answer/Resolve

1. Trigger	2. Attempt Self-Service	3. Wait/Escalate	4. Get Answer/Resolve
EXPECTATION Wants quick clarity on what happened and whether it's covered	**EXPECTATION** Hopes for a fast, specific answer	**EXPECTATION** Expects to eventually reach someone who can help	**EXPECTATION** Wants a clear answer without repeating their situation
CHALLENGE Unsure where to get an authoritative answer *[INFERRED]*	**CHALLENGE** Generic bots don't know their account, so no real answer comes back	**CHALLENGE** Long hold times	**CHALLENGE** Has to re-explain to whichever agent finally answers
OPPORTUNITY Proactive, context-aware notification could pre-empt the need to ask at all *[INFERRED]*	**OPPORTUNITY** Companion's personalization differentiator applies directly here	**OPPORTUNITY** Containment here directly serves the 50% no-escalation target	**OPPORTUNITY** Repetition is eliminated only if a future warm-transfer/context feature ships — not yet committed, flagged for Stage 9/10

EMOTIONAL ARC

Mild uncertainty at Trigger *[INFERRED]* → frustration builds through failed self-service and hold time (source-grounded) → resolution feeling depends on whether repetition is required *[INFERRED]*.

Validation status: hypothesis-based (inherited from Stage 2/3).

Provider — the time-pressed verifier

Current-state · stages: Trigger → Attempt Self-Service → Wait/Escalate → Get Answer/Resolve

1. Trigger	2. Attempt Self-Service	3. Wait/Escalate	4. Get Answer/Resolve
EXPECTATION Needs coverage/eligibility confirmed fast, often mid-workflow	**EXPECTATION** Wants a direct, actionable answer	**EXPECTATION** Minimal wait, given the time-sensitivity	**EXPECTATION** Wants an answer immediately usable, not a callback
CHALLENGE Query may involve a patient who isn't present	**CHALLENGE** Portal today isn't fast/ complete enough, so they call anyway *[INFERRED]*	**CHALLENGE** Hold time directly blocks clinical/office workflow —	**CHALLENGE** No good "call back later" path if unresolved

— a case self-service historically underserves	**OPPORTUNITY** Terse, high-density automated responses matched to provider urgency	costs more than a patient's wait	**OPPORTUNITY** Warm transfer with context (once committed) reduces follow-up need on escalation
OPPORTUNITY Companion could explicitly support patient-absent, provider-initiated lookups		**OPPORTUNITY** Faster containment on this channel has outsized operational value	

EMOTIONAL ARC

Time-pressure present from Trigger onward *[INFERRED]*, staying elevated through the workflow-blocking wait, easing only once an actionable answer arrives *[INFERRED]*.

Validation status: hypothesis-based (inherited from Stage 2/3).

Support agent — the escalation receiver (internal)

Current-state · stages: Receive Escalation → Reconstruct Context → Resolve → Close

1. Receive Escalation	2. Reconstruct Context	3. Resolve	4. Close
EXPECTATION Hopes context arrives automatically with the case	**EXPECTATION** Wants to avoid re-asking what the member already said	**EXPECTATION** Resolve accurately once context is established	**EXPECTATION** Case closed correctly, member satisfied *[INFERRED: reasonable extension]*
CHALLENGE No context is passed on transfer today	**CHALLENGE** Has to ask the member to repeat themselves, wasting time *[INFERRED]*	**CHALLENGE** Risk of acting on a noisy or wrong bot handoff	**CHALLENGE** Gap: source personas/ empathy map don't cover this stage directly
OPPORTUNITY Structured handoff summary (per Stage 3 design implication) could deliver this	**OPPORTUNITY** Automated context transfer removes this step entirely	**OPPORTUNITY** Trustworthy context transfer speeds resolution and lowers AHT	**OPPORTUNITY** Gap: not yet addressed — candidate for Stage 6 metric definition (e.g., resolution accuracy on escalations)

EMOTIONAL ARC

Skepticism at Receive *[INFERRED]* → mild frustration during Reconstruct (source-grounded) → neutral through Resolve/Close *[INFERRED]*.

Validation status: hypothesis-based (inherited from Stage 2/3).

PHASE: DIRECTION — UNCHANGED

Stage 5 — Business Model Canvas

Internal business case. Revenue Streams reframed as cost avoidance. Only Patients and Providers as Customer Segments.

Companion — Business Model Canvas

Internal business case for leadership investment · Osterwalder 9-block · hypothesis-based where noted

KEY PARTNERS

- Internal UHG/Optum systems owners (claims, benefits, provider directory) — required for integration
- Compliance/legal (HIPAA oversight for personalization using member data)
- *Open question: build vs. buy for underlying AI/NLP technology — no vendor specified in input*

KEY ACTIVITIES

- Build healthcare-specific NLP/ intelligence layer
- Integrate with claims, benefits, and provider-directory systems
- Build HIPAA-compliant personalization pipeline
- Build warm-transfer context handoff to agents
- Operate continuous-learning feedback loop
- Deliver and validate MVP in SoCal by Feb 2025

KEY RESOURCES

- Member data: claims, benefits, provider directory
- Existing UHG systems/APIs
- HIPAA compliance infrastructure and expertise
- AI/NLP capability (build or acquired)

VALUE PROPOSITIONS

Patients

- Instant, personalized answers on benefits/coverage/claims/billing without holding
- Visible use of real account data builds trust vs. generic bots
- End-to-end self-service for routine tasks

Providers

- Fast, actionable coverage/eligibility verification, including patient-absent lookups
- Minimizes disruption to clinical/office workflow vs. calling and holding

COST STRUCTURE

- AI/NLP development and ongoing maintenance
- Systems-integration engineering (claims/benefits/provider-directory APIs)
- HIPAA compliance and security infrastructure
- Continuous-learning/model-training operations
- MVP rollout costs against the Feb 2025 SoCal deadline
- *No specific dollar figures given for these cost drivers — only the $57M baseline (current agent-cost spend being reduced) was provided*

RATIONALE NOTE

This is framed as an internal investment case rather than an external product: value is captured through reduced agent-driven contact costs (Discovery's $57M baseline), not through pricing paid by patients or providers. The two served segments — patients and providers — have distinct value propositions and channels because their urgency and workflow context differ (Stage 3/4 findings), but both share the same underlying resources (member data, systems integration, compliance infrastructure) and activities. Support-team benefits (Stage 3/4) inform Key Activities and the AHT-driving cost logic, but per your direction the support team itself isn't listed as a served customer segment.

PHASE: DIRECTION — UNCHANGED

Stage 6 — Goals & Measurement (North Star)

North Star: **% of interactions resolved to the user's satisfaction without escalation.** Four pillars: Answer Trust/Accuracy, Effort/Speed, Escalation Quality, Query Coverage — all net-new instrumentation.

Companion — Goals & Measurement (North Star Framework)

First North Star for this initiative · audience: exec/stakeholder alignment · all metrics net-new instrumentation

1. Framework Overview

North Star chosen: **% of interactions resolved to the user's satisfaction without escalation.** This was chosen over "$ saved" (a lagging business outcome, not a user behavior) and over raw containment rate alone (which rewards deflection even without a genuinely good outcome). Combining resolution and satisfaction keeps the team accountable to real value delivered, not just fewer agent transfers.

Two time horizons apply, shown separately below: **MVP-era** (SoCal, live by Feb 2025 — proving the 50% containment / CSAT targets) and **steady-state** (2026 — implied by the $30M/year savings target, which requires broader rollout beyond SoCal).

2. North Star Metric

Definition: the share of Companion-initiated interactions where the user's need is fully addressed without transfer to a live agent, *and* the user's satisfaction with that resolution is confirmed (via CSAT signal or equivalent) — not resolution alone, and not satisfaction alone.

Baseline: not yet measured. No current instrumentation exists for this metric today.

3. Supporting Pillars

[ASSUMPTION: pillars proposed by this analysis, not given directly in the original input — confirmed by user before drafting]

- **Answer Trust/Accuracy** — personalization accuracy rate
  NET-NEW
  ; CSAT specific to bot-handled interactions
  NET-NEW
- **Effort/Speed** — time-to-resolution for self-service interactions
  NET-NEW
  ; repeat-contact rate for the same issue
  NET-NEW
- **Escalation Quality** — AHT delta on escalations with vs. without context passed
  NET-NEW
  ; agent-reported context usefulness
  NET-NEW
- **Query Coverage** — % of inbound volume Companion can attempt, by category
  NET-NEW
  ; containment rate by query category
  NET-NEW

4. Instrumentation Gaps

Every sub-metric above needs new tracking — none exist today (per your confirmation). Specifically missing: a CSAT survey mechanism tied to bot interactions, a personalization-accuracy auditing process, an agent-facing context-usefulness survey, and query-category tagging on inbound volume. This last gap also blocks resolving the open volume-mix question flagged back in Stage 1 — Query Coverage can't be baselined until that taxonomy exists.

5. Risks & Vanity-Metric Traps

- Raw containment % alone could be gamed — a bot that deflects without truly resolving would inflate the 50% number without real value. This is why the North Star requires satisfaction, not just non-escalation.
- CSAT collected only from users who complete a survey tends to skew positive (self-selection bias) — worth flagging before this becomes the primary quality signal.
- A drop in AHT could look like a win even if it's caused by agents skipping proper review of insufficient bot-passed context — AHT should be read alongside agent-reported context usefulness, not alone.

6. Assumptions Made Due to Missing Information

- No baseline exists for any proposed metric — all net-new, as confirmed.
- Pillars (trust/accuracy, effort/speed, escalation quality, query coverage) were proposed here, not supplied in the original brief — confirmed via "proceed."
- Query Coverage pillar is directly limited by the still-unresolved volume-mix/query-taxonomy question from Stage 1.

7. North Star Metric Tree

% of interactions resolved to the user's satisfaction without escalation

MVP-era target: 50% containment, higher CSAT (SoCal, Feb 2025) | Steady-state: scaled to support $30M/yr savings (2026)

Answer Trust / Accuracy	Effort / Speed	Escalation Quality	Query Coverage
Personalization accuracy rate NET-NEW	Time-to-resolution NET-NEW	AHT delta w/ vs. w/o context NET-NEW	% volume attemptable by category NET-NEW
Bot-interaction CSAT NET-NEW	Repeat-contact rate NET-NEW	Agent-reported context usefulness NET-NEW	Containment rate by category NET-NEW

PHASE: DIRECTION — UNCHANGED

Stage 7 — Strategic Alternatives Evaluation

Retrospective framing, scored honestly. Companion beats IVR/Staffing/Vendor alternatives on Differentiation — the one criterion tied to the actual goal — even though it loses on raw High-rating count.

Strategic Alternatives — Comparison Matrix

Simple relative scoring (High/Medium/Low) · retrospective framing · Companion is the committed direction, scored identically to the rest

CRITERION	COMPANION (COMMITTED)	IMPROVED IVR/ ROUTING	STAFFING/ PROCESS REDESIGN	VENDOR CHATBOT
Cost to implement	**Low** Custom NLP + deep integration + compliance infra + continuous learning ops (Stage 5 Cost Structure) — highest build cost of the four	**High** Extends existing IVR logic — cheapest option	**Medium** No build cost, but ongoing labor cost that doesn't shrink the $57M baseline it's meant to fix	**Medium** Licensing fees plus still-required integration work
Time-to-market	**Medium** Feasible for SoCal MVP scope, but integration readiness with claims/benefits/ provider-directory systems is still an open risk (Stage 1)	**High** Fastest — incremental change to infrastructure already in place	**High** No engineering build required, can move on org/process timelines	**Medium** Faster than custom build, but still needs integration and compliance validation before launch
Differentiation / HIPAA-specific personalization fit	**High** Purpose-built for HIPAA-compliant personalization and deep UHG systems integration — the only option that can genuinely resolve routine queries, not just route them faster	**Low** Still ends at a live agent — doesn't reduce agent-handled volume, the actual problem from Discovery	**Low** No technology differentiation, doesn't scale, doesn't add self-service capability	**Low** Generic solution unlikely to match healthcare-specific, deeply-integrated personalization out of the box
Operational feasibility	**Medium** Requires sustained coordination with claims/ benefits/provider-directory system owners (Stage 5 Key Partners)	**High** Minimal new systems, easiest to operate	**Medium** Depends on hiring/ process execution and change management	**Medium** Vendor dependency reduces control over roadmap and the continuous-learning loop
Regulatory/ compliance risk	**Medium** Directly handles PHI for personalization — elevated but internally controlled and audited	**Low risk** No new PHI-handling channel	**Low risk** No new tech or data exposure beyond current state	**High risk** Third-party handling of PHI raises BAA/ compliance exposure, harder to audit than an internal build
High-rating count (unweighted, informational only — not a decision rule)	**1 of 5**	**4 of 5**	**2 of 5**	**0 of 5**

PHASE: STRATEGY — CHANGED: THEME LEVEL

Stage 8 — Prioritization

Per the updated skill's granularity check, candidates are now 5 strategic themes instead of 10 already-atomic features.

THEME	ROLLS UP	IMPACT	EFFORT
A. Personalized Self-Service	Personalized answer engine, claims integration, benefits integration, query taxonomy	High	High
B. Agent Handoff Quality	Warm transfer with context	High	Medium
C. Provider Support Expansion	Provider directory, patient-absent lookup, tone differentiation	Medium	Medium-High
D. Continuous Learning & Model Improvement	Continuous learning loop	Medium	Medium
E. Outcome Measurement	CSAT instrumentation	High	Low

Theme-Level Impact vs. Effort

5 themes, replacing the original 10 feature-level candidates

A Personalized Self-Service **B** Agent Handoff Quality

C Provider Support Expansion **D** Continuous Learning

E Outcome Measurement

Recommended Sequencing

E (Quick Win) → A, B (Major Projects, unavoidable core) → C, D (deferred).

PHASE: STRATEGY — CHANGED: THEME LEVEL

Stage 9 — MVP Scope

Hard constraint: live in SoCal by Feb 2025. MVP floor still qualitative: [ASSUMPTION] measurable, positive containment and CSAT signal above the 0% baseline — no specific percentage confirmed. Same substantive conclusion as v1, now expressed at theme level.

MVP Scope Boundary — Theme Level (Updated)

Re-derived from the Stage 8 theme-level re-run — same substantive conclusion as the original feature-level pass

IN — MVP (FEB 2025, SOCAL)

A. Personalized Self-Service

Stage 8: High Impact / High Effort

Rolls up: personalized answer engine, claims integration, benefits integration

B. Agent Handoff Quality

Stage 8: High Impact / Medium Effort

Rolls up: warm transfer with context

E. Outcome Measurement

Stage 8: High Impact / Low Effort

Rolls up: CSAT instrumentation

OUT — DEFERRED

C. Provider Support Expansion

Stage 8: Medium Impact / Med-High Effort

Rolls up: provider directory, patient-absent lookup, tone differentiation

D. Continuous Learning & Model Improvement

Stage 8: Medium Impact / Medium Effort

Rolls up: continuous learning loop

Still load-bearing: MVP is patient-first. Never explicitly confirmed by leadership — carried forward unchanged from v1.

PHASE: STRATEGY — CHANGED: TECHNICAL CONSIDERATIONS ADDED

Stage 10 — PRD (Companion, SoCal MVP)

Problem Statement

Optum routes nearly all patient and provider questions — including simple, repetitive ones — to live agents, costing ~$57M/year and producing long waits and poor UX.

Goal & Success Metrics

[ASSUMPTION] no specific MVP-floor number confirmed — positive containment and CSAT above the 0% baseline, live in SoCal.

Scope

In (theme level): A. Personalized Self-Service, B. Agent Handoff Quality, E. Outcome Measurement. **Out:** C. Provider Support Expansion, D. Continuous Learning & Model Improvement.

Technical Considerations (new section)

Did not exist in v1 — added per the updated skill. Proposed NFR defaults and integration architecture below, all flagged as recommendations pending real sign-off, plus a rollout/rollback plan.

NFR	STATUS
Response latency	No target given. [RECOMMENDATION] sub-3-second first response, needs engineering confirmation.
Uptime/availability	No target given. [RECOMMENDATION] match or exceed current IVR/portal uptime — that baseline itself wasn't provided either.
Throughput/peak concurrency	Can't be sized — blocked by the same unresolved Stage 1 volume-mix gap.
Data retention (audit logs)	HIPAA baseline: 6 years (45 CFR §164.316(b)(2)(i)) — general rule, specific policy still needs compliance sign-off.
Data residency	Not addressed anywhere — open question.

Proposed default, not confirmed by any systems team — flagged for sign-off, not settled architecture.

Companion — Proposed System Architecture (Draft)

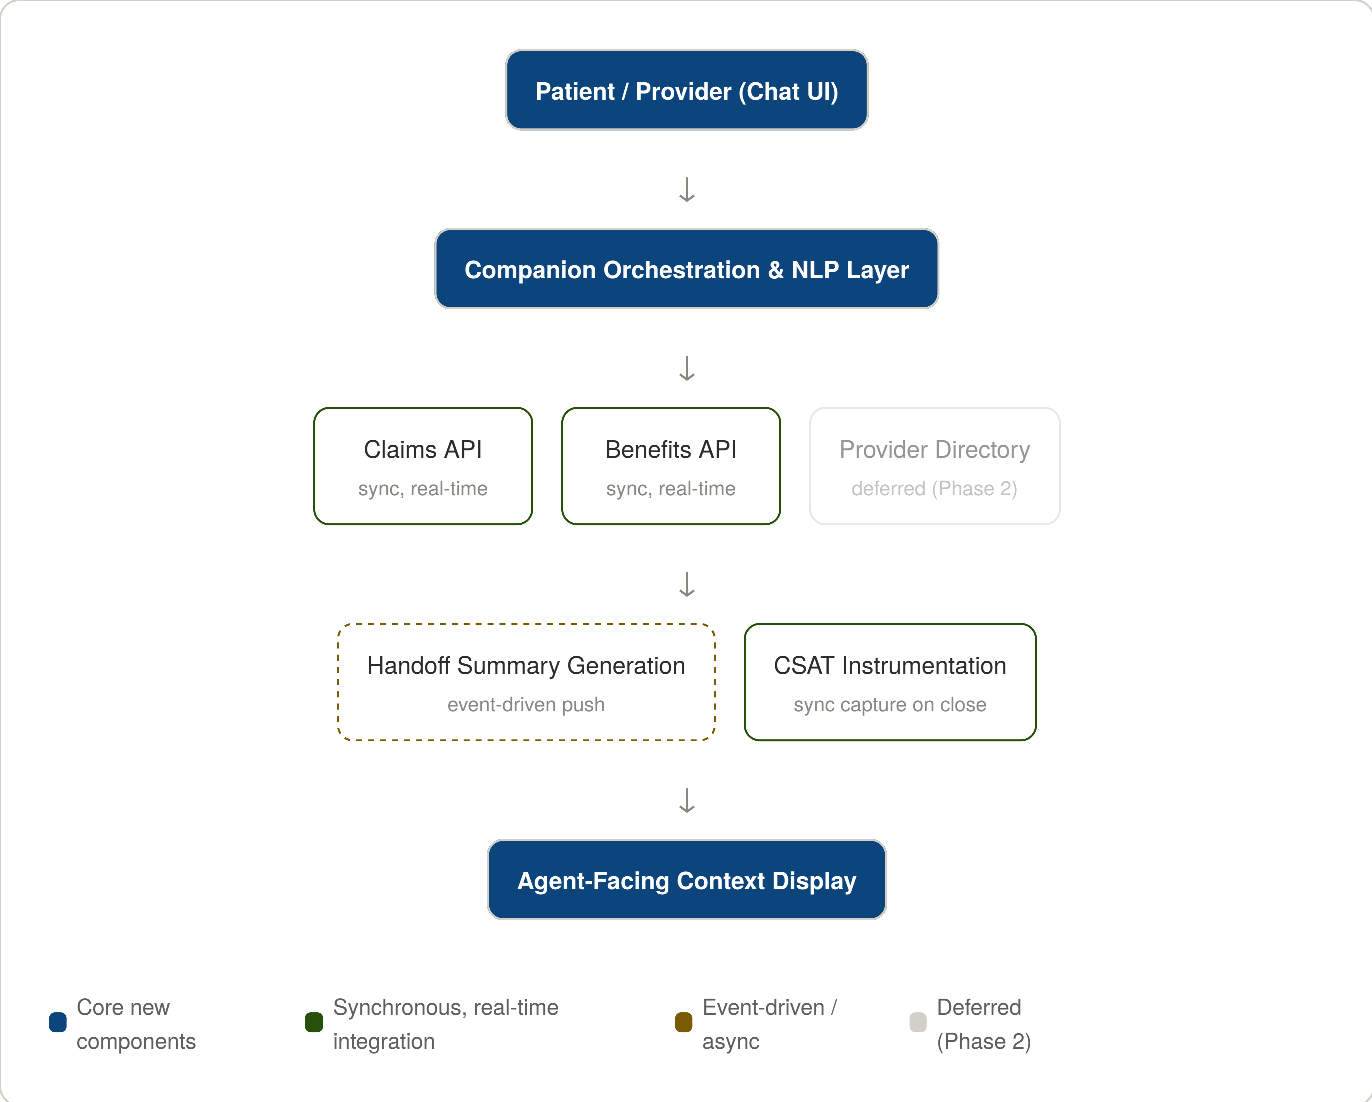

Rollout/Rollback [RECOMMENDATION, needs confirmation]: staged rollout to a subset of eligible SoCal members before 100%; an explicit kill switch to revert fully to today's agent-routing baseline; rollback triggers tied to CSAT/containment falling materially below the (still-unconfirmed) MVP floor, or any audit/compliance failure.

Risks & Open Questions

- Claims/benefits integration readiness is unconfirmed.
- No specific MVP-floor number exists.
- Patient-first assumption means providers get zero MVP improvement.
- No NFR numbers are actually confirmed — every value above is a proposed default, not a given fact.

PHASE: STRATEGY — CHANGED: SIMPLIFIED TO 3 THEMES

Stage 11 — Roadmap

Simplified from v1's 5-in/5-out feature split to 3 themes in Phase 1, 1 in Phase 2, 1 in Phase 3. Illustrative dates retained per earlier confirmation.

Companion — Roadmap (Theme-Level, Updated)

Simplified from the original 5-in/5-out feature split to 3 themes in, 2 deferred

Feb 2025 / 2026 shown as illustrative dates, per earlier confirmation

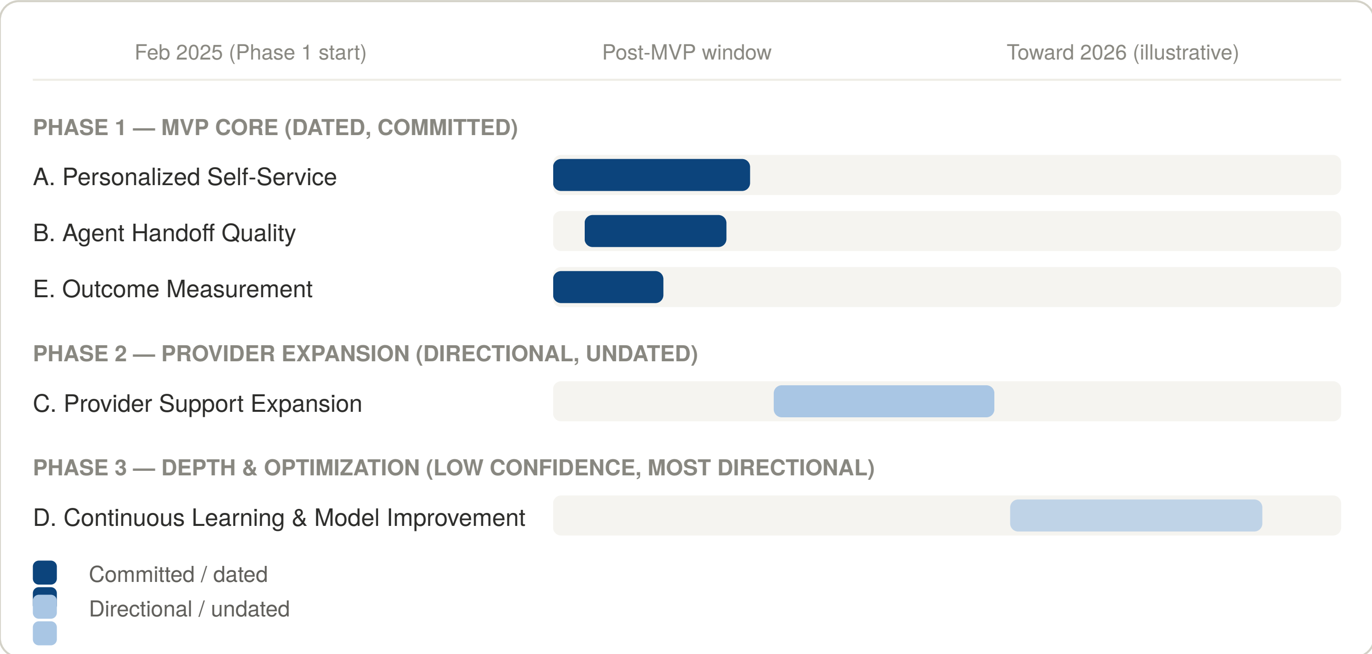

Phase gates unchanged in substance: Gate 1→2 requires the MVP floor hit, integrations stable, warm transfer measurably reducing AHT. Gate 2→3 requires provider directory launched/adopted and sufficient live data.

PHASE: EXECUTION — UNCHANGED

Stage 12 — GTM Plan

No change from v1 — the patient-facing value prop and channel strategy were already about the outcome, not internal feature/theme naming, so this doesn't depend on the granularity shift.

Companion — GTM Visuals (Phase 1, Patients)

Single segment in scope for this plan — providers deferred to a future GTM plan alongside Roadmap Phase 2

Positioning & Channel Matrix

SEGMENT	BEHAVIOR DISPLACED	CORE MESSAGE	CHANNEL
Patients	Calling and holding, or using a generic portal that doesn't reflect real account data	"Get real answers about your coverage, claims, and bills instantly — no hold time, no re-explaining yourself."	Email (pending claims/bills)
			In-app notification/ banner at point of need
			IVR offer before entering the hold queue

Launch Sequencing Timeline

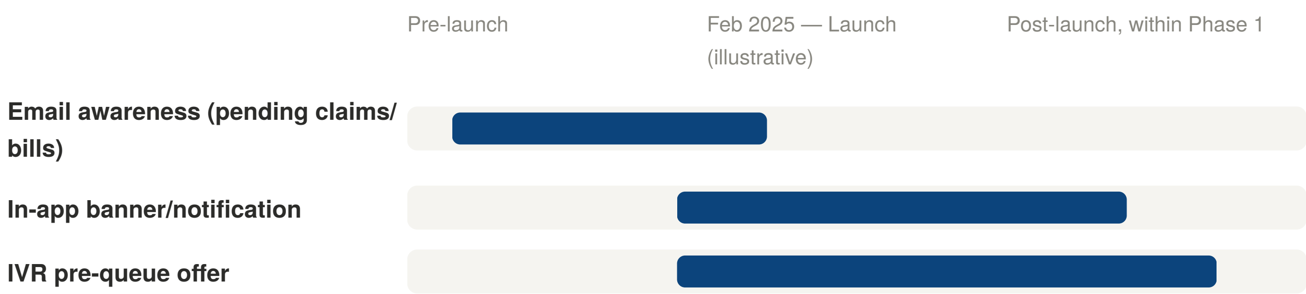

Positioned within the roadmap's Phase 1 window (Feb 2025, illustrative) — no separate GTM-only phase structure introduced. Relative sequencing only; no day-level schedule was given.

PHASE: EXECUTION — CHANGED: 6 FEATURES, 1 GENUINE SPLIT

Stage 13 — Feature Detail

Decomposing Theme A converges on the same three features as v1 (unchanged): **Personalized Answer Engine, Claims Integration, Benefits Integration**. Theme E converges on the same single feature: **CSAT Instrumentation**.

Theme B decomposed differently — the genuine new finding the theme-first approach enabled. The original "Warm Transfer with Context" bundled a backend concern and a frontend concern that don't share an owner:

Handoff Summary Generation (new split)

Problem: When Companion can't resolve a request, there's no structured record of what was asked/found to hand off.

Requirements: Generate a structured summary of the request and any data surfaced, on escalation.

Owner: Backend/NLP.

Agent-Facing Context Display (new split)

Problem: Even with a summary generated, the agent has no interface to see it before the call connects.

Requirements: Display the handoff summary to the assigned agent, gated by access control.

Owner: Frontend/agent tooling.

Open question: The actual agent-facing UI still isn't defined — same gap as v1, now correctly scoped to the feature that actually owns it.

Personalized Answer Engine, Claims Integration, Benefits Integration, and CSAT Instrumentation are unchanged from v1 — full detail blocks not repeated here.

PHASE: EXECUTION — CHANGED: RICE-RANKED

Stage 14 — User Stories

16 stories across the now-6 features (the Theme B split added 4 stories where v1 had 3). Ranked by **RICE** (Reach × Impact × Confidence ÷ Effort) instead of v1's Status-based Ready/Blocked ordering.

Reach/Impact/Confidence/Effort are relative 1–5 judgment scores (Confidence uses standard 0.5/0.8/1.0 bands) — illustrative estimates, not measured telemetry. Ranked by RICE score, descending.

Stage 14 Backlog — RICE-Ranked

16 stories across 6 features (post theme-level re-run)

STORY	FEATURE	REACH	IMPACT	CONF.	EFFORT	RICE
CSAT-1 Capture CSAT Signal	CSAT Instrumentation	5	4	0.8	2	8.00
PAE-3 Identity Verification	Personalized Answer Engine	5	4	0.8	3	5.33
BI-1 Real-Time Benefits Lookup	Benefits Integration	4	5	0.8	3	5.33
PAE-1 Personalized Response	Personalized Answer Engine	5	5	0.8	5	4.00
PAE-4 Audit Logging (data access)	Personalized Answer Engine	5	2	0.8	2	4.00
CI-1 Real-Time Claim Status Lookup	Claims Integration	4	5	0.8	4	4.00
HSG-2 Audit Logging (summary gen)	Handoff Summary Generation	3	2	0.5	1	3.00
HSG-1 Structured Summary Generation	Handoff Summary Generation	3	5	0.5	3	2.50
ACD-2 Access Control on Summary View	Agent-Facing Context Display	3	3	0.5	2	2.25
ACD-1 Display Summary to Agent	Agent-Facing Context Display	3	4	0.5	3	2.00
CSAT-2 Fallback: No CSAT Response	CSAT Instrumentation	2	2	0.5	1	2.00
PAE-2 Fallback: Data Unavailable	Personalized Answer Engine	2	3	0.5	2	1.50
CI-3 Scope Boundary (provider/absent)	Claims Integration	1	2	0.5	1	1.00
BI-3 Scope Boundary (provider/absent)	Benefits Integration	1	2	0.5	1	1.00
CI-2 Fallback: Claims System Down	Claims Integration	1	3	0.5	2	0.75
BI-2 Fallback: Benefits System Down	Benefits Integration	1	3	0.5	2	0.75

What RICE Changed vs. v1's Status Ordering

CSAT-1 (capture CSAT signal) jumps to #1 — near-total reach, low effort, and it's the only way to know if anything else worked. Fallback/scope-boundary stories sink to the bottom even though their parent features are otherwise high-priority.

Reach/Impact/Confidence/Effort are relative judgment scores, not measured telemetry — nothing has shipped yet.

PHASE: VALIDATE — VERDICT UNCHANGED, SCORECARD ADDED

Stage 15 — PMF Check

Assessment, Patients/SoCal, threshold = MVP floor/North Star metrics. Evidence: none — pre-launch.

Companion — PMF Signal Scorecard

Assessment · Patients, SoCal · pre-launch — zero evidence exists

SIGNAL	RATING	EVIDENCE
Retention / repeat usage	**No data**	Pre-launch, nothing to measure
Organic / referral growth	**No data**	Pre-launch
Engagement intensity	**No data**	Pre-launch
Qualitative sentiment	**No data**	Pre-launch
Containment rate *(chosen threshold)*	**No data**	Instrumentation defined, not yet live
CSAT *(chosen threshold)*	**No data**	Instrumentation defined, not yet live

Evidence Chart

No data available

Pre-launch — no retention, usage, or cohort data exists to plot.
Populate this chart once containment/CSAT instrumentation goes live post-Feb 2025 launch.

Real Verdict

Insufficient evidence to call this. Unchanged from v1 — the theme-level restructuring of Stages 8–14 doesn't create evidence that doesn't exist. Get the MVP live, then reassess with real containment/CSAT/retention data.

Illustrative Scorecard (Best-Judgment Ratings)

Hypothetical numbers only, requested separately to demonstrate the analysis method. Ratings are reasoned from actual pipeline findings — not random — and don't change the real verdict above.

⚠ ILLUSTRATIVE — best-judgment estimates, not real evidence. Companion is pre-launch. The real Stage 15 verdict remains "insufficient evidence." Ratings are reasoned from actual pipeline findings (design attention, GTM channel strategy, unresolved data gaps), not arbitrary.

Companion — PMF Signal Scorecard (Best-Judgment Ratings)

Hypothetical ~90-day post-launch snapshot · Patients, SoCal

NORTH STAR PILLARS (STAGE 6)

PILLAR	ILLUSTRATIVE VALUE	RATING	WHY (JUDGMENT BASIS)
Answer Trust / Accuracy	Personalization accuracy 91% · Bot CSAT 76%	**Strong**	Maps to the most-scrutinized capability across Stages 2–4 (personalization was the core differentiator); reasonable to judge it performs near spec.
Effort / Speed	Avg. resolution 2.5 min · Repeat-contact rate 22%	**Mixed**	Speed likely fine; a 22% repeat-contact rate suggests real cases still aren't fully resolved first-touch.
Escalation Quality	AHT reduced ~18% · 68% agents rate context useful	**Emerging**	Warm transfer got heavy design attention (Stage 3's top design implication), but 68% agent approval leaves real room, and the agent-facing UI was never fully specified (Stage 13).
Query Coverage	~45% of inbound volume attemptable	**Weak**	Direct consequence of the still-unresolved Stage 1 volume-mix gap and the deferred query-taxonomy work — this is the pillar most predictably weak by design, not by chance.

BROADER PMF SIGNALS

SIGNAL	ILLUSTRATIVE VALUE	RATING	WHY (JUDGMENT BASIS)
Containment rate (MVP floor)	38%	**Emerging**	Below the eventual 50% target but a meaningful non-zero floor, consistent with Query Coverage capping what's attemptable.
30-day retention	21% return usage	**Weak**	Nothing in the product or GTM plan specifically drives repeat usage — a single resolved query doesn't create an obvious reason to return.
Organic / referral growth	15% of usage	**Weak**	GTM (Stage 12) is entirely push-driven — email, IVR, in-app — with no organic mechanism designed in. Low organic share is the predictable result, not a surprise.

SIGNAL CHART

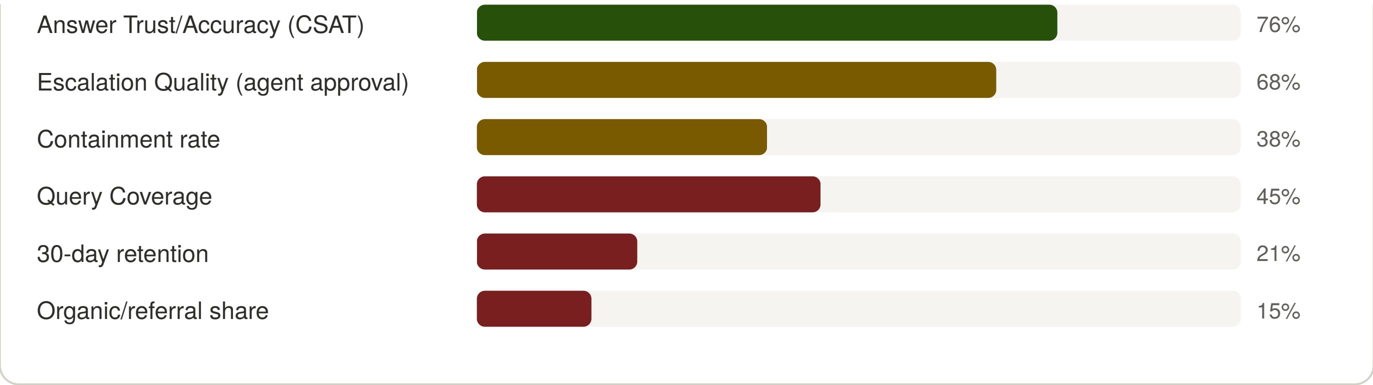

Overall illustrative read: Mixed-but-promising. The weak points (Query Coverage, organic growth, retention) trace directly to specific unresolved pipeline decisions: the Stage 1 volume-mix gap, the push-only GTM channel strategy, and the absence of anything designed to drive return visits — not generic PMF caution.

---

v3 of this report reflects the updated pm-discovery-to-execution skill: theme-level Prioritization (8) cascading through MVP Scope (9), PRD (10, now with Technical Considerations), Roadmap (11), and Feature Detail (13, where the theme-first approach surfaced a genuine feature split v1 had missed) — plus RICE-ranked backlog ordering in User Stories (14) and an illustrative PMF scorecard in Stage 15. Unresolved threads carried across all versions: the patient-first MVP assumption was never independently confirmed by leadership; the volume-mix/query-taxonomy data gap from Stage 1 was never resolved; no specific MVP-floor number or NFR target exists anywhere in the pipeline.